

# Lecture 2: Problem as Search in a Graph

**CSSE 5600/6600: Artificial Intelligence**

**Instructor: Bo Liu**

# Content

- ❖ Why is search the key problem-solving technique in AI?
- ❖ Problem types
- ❖ Problem formulation
- ❖ Understanding and comparing several “blind” search algorithms.

# Search examples

- Water jugs: how to get 1?




7



5

- Two operations: fill up (from tap or other jug), empty (to ground or other jug)




Rules:

- (1) boat can take at most 2 objects one time
- (2) F kills D, D kills S, S kills C, and C kills F, if no others around.

Goal: Everybody go to the other side of this river

# The search problem

- **State space**  $S$  : all valid configurations
- **Initial states (nodes)**  $I = \{(CSDF, )\} \subseteq S$ 
  - Where's the boat?
- **Goal states**  $G = \{(\cdot, CSDF)\} \subseteq S$
- **Successor function**  $succs(s) \subseteq S$  : states reachable in one step (one arc) from  $s$ 
  - $succs((CSDF, )) = \{(CD, SF)\}$
  - $succs((CDF, S)) = \{(CD, FS), (D, CFS), (C, DFS)\}$
- **Cost**  $(s, s') = 1$  for all arcs. (weighted later)
- The search problem: find a solution path from a state in  $I$  to a state in  $G$ .
  - Optionally minimize the cost of the solution.



# Search examples

- 8-puzzle

7	2	4
5		6
8	3	1

Start State


	1	2
3	4	5
6	7	8

Goal State

- States = configurations
- successor function = up to 4 kinds of movement
- Cost = 1 for each move

# Search examples

- Water jugs: how to get 1?



- Goal? (How many goal states?  $(0,0), (0,1), \dots, (7,5)$ )
- Successor functions: fill up (from tap or other jug), empty (to ground or other jug) – fill big jug, fill small jug, empty big jug, empty small jug, fill big jug from small jug, fill small jug from small jug

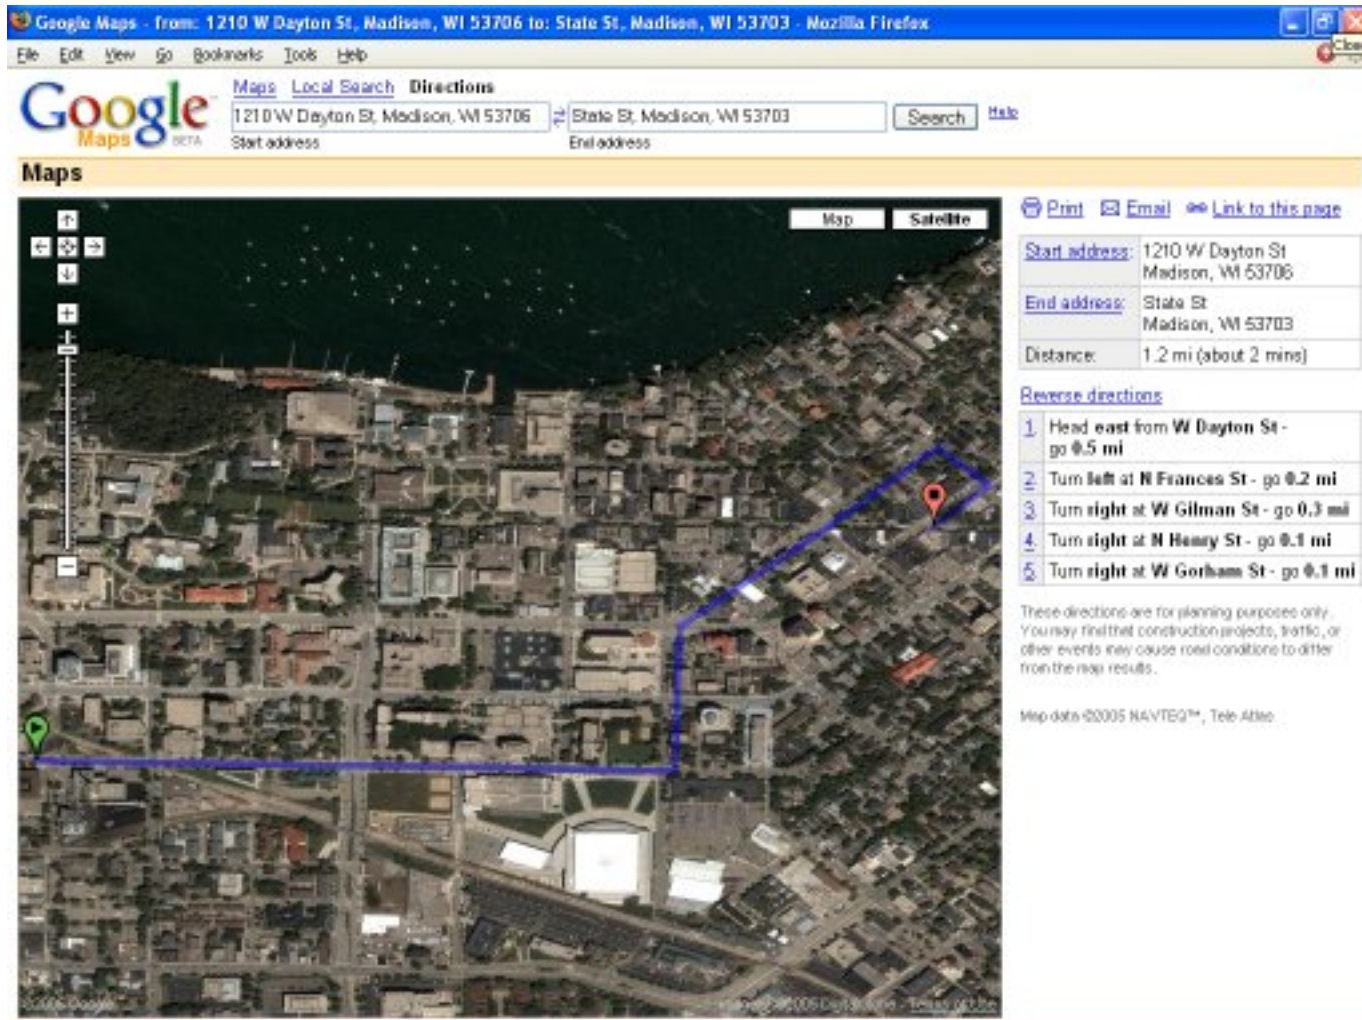
## One solution

- $(7, 0) \rightarrow (2, 5) \rightarrow (2, 0) \rightarrow (2, 0) \rightarrow (7, 2) \rightarrow (4, 5) \rightarrow (4, 0) \rightarrow (0, 4) \rightarrow (7, 4) \rightarrow (6, 5) \rightarrow (6, 0) \rightarrow (1, 5) \rightarrow (1, 0)$
- Other solutions?




# Search examples

- Route finding (state? Successors? Cost weighted)



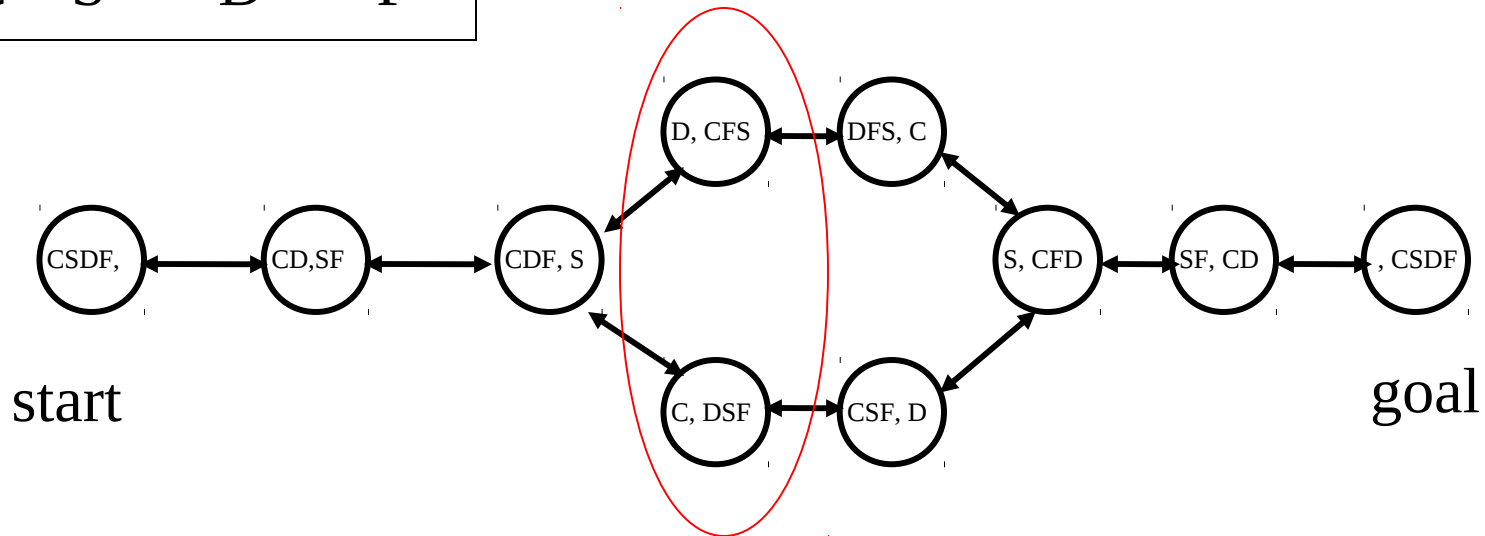

# 8-queens

- How to define states?



- states? -any arrangement of  $n \leq 8$  queens  
-or arrangements of  $n \leq 8$  queens in leftmost  $n$  columns, 1 per column, such that no queen attacks any other.
- initial state? no queens on the board
- Successor functions? -add queen to any empty square  
-or add queen to leftmost empty square such that it is not attacked by other queens.
- goal test? 8 queens on the board, none attacked.
- path cost? 1 per move


# A directed graph in state space



- In general there will be many generated, but un-expanded states at any given time
- One has to choose which one to expand next

# Different search strategies

- The generated, but not yet expanded states form the **fringe (OPEN)**.
- The essential difference is **which one to expand first**.
- Deep or shallow?



# Uninformed search on trees


- **Uninformed** means we only know:
  - The goal test
  - The *succs()* function
- But **not** which non-goal states are better: that would be informed search (next lecture).
- For now, we also assume *succs()* graph is **a tree**.
  - Won't encounter repeated states.
  - We will relax it later.
- Search strategies: BFS, UCS, DFS, IDS, BIBFS
- Differ by what un-expanded nodes to expand

# Breadth-first search (BFS)

Expand the shallowest node first

- Examine states **one** step away from the initial states
- Examine states **two** steps away from the initial states
- and so on...


ripple



# Breadth-first search (BFS)

Use a **queue** (First-in First-out)

1. en\_queue(Initial states)
2. While (queue not empty)
3.     s = de\_queue()
4.     if (s==goal) success!
5.     T = succs(s)
6.     en\_queue(T)
7. endwhile






# Breadth-first search (BFS)

Use a **queue** (First-in First-out)

1. en\_queue(Initial states)
2. While (queue not empty)
3.     s = de\_queue()
4.     if (s==goal) success!
5.     T = succs(s)
6.     en\_queue(T)
7. endwhile




queue (**fringe**, **OPEN**)  
→ [A] →

# Breadth-first search (BFS)

Use a **queue** (First-in First-out)

1. en\_queue(Initial states)
2. While (queue not empty)
3.     s = de\_queue()
4.     if (s==goal) success!
5.     T = succs(s)
6.     en\_queue(T)
7. endwhile




queue (**fringe**, **OPEN**)  
→ [CB] → A

# Breadth-first search (BFS)

Use a **queue** (First-in First-out)

1. en\_queue(Initial states)
2. While (queue not empty)
3.   s = de\_queue()
4.   if (s==goal) success!
5.   T = succs(s)
6.   en\_queue(T)
7. endwhile




queue (**fringe**, **OPEN**)  
→ [EDC] → B

# Breadth-first search (BFS)

Use a **queue** (First-in First-out)

1. en\_queue(Initial states)
2. While (queue not empty)
3.     s = de\_queue()
4.     if (s==goal) success!
5.     T = succs(s)
6.     en\_queue(T)
7. endwhile




queue (**fringe** , **OPEN**)  
→[GFED] → C

If G is a goal, we've seen it, but we don't stop!

# Breadth-first search (BFS)

Use a **queue** (First-in First-out)

- en\_queue(Initial states)
- While (queue not empty)
- s = de\_queue()
- if (s==goal) success!
- T = succs(s)
- **for t in T: t.prev=s**
- en\_queue(T)
- endwhile



queue  
→ [] → G


... until much later we pop G.

We need **back pointers** to recover the solution path.

Looking stupid?  
Indeed. But let's be  
consistent...

# Performance of BFS

- Assume:
  - the graph may be infinite.
  - Goal(s) exists and is only finite steps away.
- Will BFS find at least one goal?
- Will BFS find the least cost goal?
- Time complexity?
  - # states generated
  - Goal  $d$  edges away
  - Branching factor  $b$
- Space complexity?
  - # states stored




# Performance of BFS

Four measures of search algorithms:

- **Completeness** (not finding all goals): yes, BFS will find a goal.
- **Optimality**: yes if edges cost 1 (more generally positive non-decreasing in depth), no otherwise.
- **Time** complexity (worst case): goal is the last node at radius  $d$ .
  - Have to generate all nodes at radius  $d$ .
  - $b + b^2 + \dots + b^d \sim O(b^d)$
- **Space** complexity (bad)
  - Back pointers for all generated nodes  $O(b^d)$
  - The queue / fringe (smaller, but still  $O(b^d)$ )

# What's in the fringe (queue) for BFS?

- Convince yourself this is  $O(b^d)$





# Performance of search algorithms on trees

b: branching factor (assume finite)    d: goal depth

	Complete	optimal	time	space
Breadth-first search	Y	Y, if <sup>1</sup>	$O(b^d)$	$O(b^d)$

1. Edge cost constant, or positive non-decreasing in depth

# Performance of BFS

Four measures of search algorithms:

**Solution:  
Uniform-cost  
search**

- **Completeness** (not finding all goals): find a goal.
- **Optimality**: yes if edges cost 1 (more generally positive non-decreasing with depth), **no otherwise**.
- **Time** complexity (worst case): goal is the last node at radius  $d$ .
  - Have to generate all nodes at radius  $d$ .
  - $b + b^2 + \dots + b^d \sim O(b^d)$
- **Space** complexity (bad, Figure 3.11)
  - Back points for all generated nodes  $O(b^d)$
  - The queue (smaller, but still  $O(b^d)$ )

# Uniform-cost search


- Find the least-cost goal
- Each node has a path cost from start (= sum of edge costs along the path). Expand the least cost node first.
- Use a **priority queue** instead of a normal queue
  - Always take out the least cost item
  - Remember *heap*? time  $O(\log(\text{\#items in heap}))$

That's it<sup>\*</sup>

<sup>\*</sup> Complications on graphs (instead of trees). Later.

# Uniform-cost search (UCS)

- Complete and optimal (if edge costs  $\geq \epsilon > 0$ )
- Time and space: can be much worse than BFS
  - Let  $C^*$  be the cost of the least-cost goal
  - $O(b^{C^*/\epsilon})$ , possibly  $C^*/\epsilon \gg d$



# Performance of search algorithms on trees

b: branching factor (assume finite)    d: goal depth

	Complete	optimal	time	space
Breadth-first search	Y	Y, if <sup>1</sup>	$O(b^d)$	$O(b^d)$
Uniform-cost search <sup>2</sup>	Y	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$

1. edge cost constant, or positive non-decreasing in depth
2. edge costs  $\geq \epsilon > 0$ .  $C^*$  is the best goal path cost.

# General State-Space Search Algorithm

```
function general-search(problem, QUEUEING-FUNCTION)
;; problem describes the start state, operators, goal test, and
;; operator costs
;; queueing-function is a comparator function that ranks two states
;; general-search returns either a goal node or "failure"

nodes = MAKE-QUEUE(MAKE-NODE(problem.INITIAL-STATE))
loop
  if EMPTY(nodes) then return "failure"
  node = REMOVE-FRONT(nodes)
  if problem.GOAL-TEST(node.STATE) succeeds
    then return node
  nodes = QUEUEING-FUNCTION(nodes, EXPAND(node,
                                              problem.OPERATORS))

;; succ(s)=EXPAND(s, OPERATORS)
;; Note: The goal test is NOT done when nodes are generated
;; Note: This algorithm does not detect loops
end
```

# Recall the bad space complexity of BFS

Four measures of search algorithms:

- **Completeness** (not finding all goals): find a goal.
- **Optimality**: yes if edges cost 1 (more generally positive non-decreasing with depth), **no otherwise**.
- **Time** complexity (goal is the last node at radius  $d$ ):
  - Have to generate all nodes at radius  $d$ .
  - $b + b^2 + \dots + b^d \sim O(b^d)$
- **Space** complexity (bad, Figure 3.11)
  - Back points for all generated nodes  $O(b^d)$
  - The queue (smaller, but still  $O(b^d)$ )

**Solution:**  
**Uniform-cost**  
**search**


**Solution:**  
**Depth-first**  
**search**

# Depth-first search

Expand the deepest node first

1. Select a direction, go deep to the end ————
2. Slightly change the end ————
3. Slightly change the end some more... ————

fan






# Depth-first search (DFS)

Use a **stack** (First-in Last-out)

1. push(Initial states)
2. While (stack not empty)
3.     s = pop()
4.     if (s==goal) success!
5.     T = succs(s)
6.     push(T)
7. endwhile



stack (**fringe**)


[] ⇔

# What's in the fringe for DFS?

- $m$  = maximum depth of graph from start


- $m(b-1) \sim O(mb)$

(Space complexity)




- “backtracking search” even less space
  - generate siblings (if applicable)

c.f. BFS  $O(b^d)$




# What's wrong with DFS?

- Infinite tree: may not find goal (incomplete)
- May not be optimal
- Finite tree: may visit almost all nodes, time complexity  $O(b^m)$



c.f. BFS  $O(b^d)$



# Performance of search algorithms on trees

b: branching factor (assume finite)    d: goal depth    m: graph depth


	Complete	optimal	time	space
Breadth-first search	Y	Y, if <sup>1</sup>	$O(b^d)$	$O(b^d)$
Uniform-cost search <sup>2</sup>	Y	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
Depth-first search	N	N	$O(b^m)$	$O(bm)$

1. edge cost constant, or positive non-decreasing in depth
  - edge costs  $\geq \epsilon > 0$ .  $C^*$  is the best goal path cost.

# How about this?

1. DFS, but stop if path length  $> 1$ .
2. If goal not found, repeat DFS, stop if path length  $> 2$ .
3. And so on...

fan within ripple



# Iterative deepening

- Search proceeds like BFS, but fringe is like DFS
  - Complete, optimal like BFS
  - Small space complexity like DFS
- A huge waste?
  - Each deepening repeats DFS from the beginning
  - No!  $db + (d-1)b^2 + (d-2)b^3 + \dots + b^d \sim O(b^d)$
  - Time complexity like BFS
- Preferred uninformed search method

# Performance of search algorithms on trees

b: branching factor (assume finite)    d: goal depth    m: graph depth

	Complete	optimal	time	space
Breadth-first search	Y	Y, if <sup>1</sup>	$O(b^d)$	$O(b^d)$
Uniform-cost search <sup>2</sup>	Y	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
Depth-first search	N	N	$O(b^m)$	$O(bm)$
Iterative deepening	Y	Y, if <sup>1</sup>	$O(b^d)$	$O(bd)$

1. edge cost constant, or positive non-decreasing in depth
  - edge costs  $\geq \epsilon > 0$ .  $C^*$  is the best goal path cost.

# Performance of search algorithms on trees

b: branching factor (assume finite)    d: goal depth    m: graph depth

	Complete	optimal	time	space
Breadth-first search	Y		$O(b^d)$	$O(b^d)$
Uniform cost search			$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
Depth-limited search			$O(b^m)$	$O(bm)$
Iterative deepening search		Y, if $b$ is finite	$O(b^d)$	$O(bd)$


How to reduce the number of states we have to generate?

1. edge cost constant, or positive non-decreasing in depth
  - edge costs  $\geq \epsilon > 0$ .  $C^*$  is the best goal path cost.




# Bidirectional search

- Breadth-first search from both start and goal
- Fringes meet
- Generates  $O(b^{d/2})$  instead of  $O(b^d)$  nodes



# Bidirectional search

- But
  - The fringes are  $O(b^{d/2})$
  - How do you start from the 8-queens goals?



# Performance of search algorithms on trees


b: branching factor (assume finite)    d: goal depth    m: graph depth

	Complete	optimal	time	space
Breadth-first search	Y	Y, if <sup>1</sup>	$O(b^d)$	$O(b^d)$
Uniform-cost search <sup>2</sup>	Y	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
Depth-first search	N	N	$O(b^m)$	$O(bm)$
Iterative deepening	Y	Y, if <sup>1</sup>	$O(b^d)$	$O(bd)$
Bidirectional search <sup>3</sup>	Y	Y, if <sup>1</sup>	$O(b^{d/2})$	$O(b^{d/2})$

1. edge cost constant, or positive non-decreasing in depth
  - edge costs  $\geq \epsilon > 0$ .  $C^*$  is the best goal path cost.
  - both directions BFS; not always feasible.

# If state space graph is not a tree

- The problem: repeated states



- Ignore the danger of repeated states: wasteful (BFS) or impossible (DFS). Can you see why?
- How to prevent it?


## If state space graph is not a tree

- We have to remember already-expanded states (**CLOSED**).
- When we take out a state from the fringe (OPEN), check whether it is in CLOSED (already expanded).
  - If yes, throw it away.
  - If no, expand it (add successors to OPEN), and move it to CLOSED.

# If state space graph is not a tree


- BFS:
  - Still  $O(b^d)$  space complexity, not worse
- DFS:
  - Known as Memorizing DFS (MEMDFS)
    - Space and time now  $O(\min(N, b^M))$  – much worse!
    - N: number of states in problem
    - M: length of longest cycle-free path from start to anywhere
  - Alternative: Path Check DFS (PCDFS): remember only expanded states on current path (from start to the current node)
    - Space  $O(M)$
    - Time  $O(b^M)$

# Example



# Example

- Depth-First Search: **S A D E G**  
Solution found: **S A G**
- Breadth-First Search: **S A B C D E G**  
Solution found: **S A G**
- Uniform-Cost Search: **S A D B C E G**  
Solution found: **S B G** (This is the only uniform-cost search algorithm that worries about costs.)
- Iterative-Deepening Search: **S A B C S A D E G**  
Solution found: **S A G**





# Depth-First Search

expanded  
node


-----

S  
A  
D  
E  
G

nodes list

-----

{ S }  
{ A B C }  
{ D E G B C }  
{ E G B C }  
{ G B C }  
{ B C }



Solution path found is S A G <-- this G has cost 10  
Number of nodes expanded (including goal node) = 5

# Breadth-First Search

expanded  
node

----

S

A

B

C

D

E

G

nodes list

-----

{ S }

{ A B C }

{ B C D E G }


{ C D E G G' }

{ D E G G' G'' }

{ E G G' G'' }

{ G G' G'' }

{ G' G'' }



Solution path found is S A G <-- this G also has cost 10  
 Number of nodes expanded (including goal node) = 7

# Uniform-Cost Search


expanded  
node

nodes list

-----

-----

	{ S }	
S	{ A(1) B(5) C(8) }	
A	{ D(4) B(5) C(8) E(8) G(10) }	(note, we don't return G)
D	{ B(5) C(8) E(8) G(10) }	
B	{ C(8) E(8) G(9) G(10) }	
C	{ E(8) G(9) G(10) G(13) }	
E	{ G(9) G(10) G(13) }	
G	{ }	



Solution path found is S B G <-- this G has cost 9, not 10  
 Number of nodes expanded (including goal node) = 7

# Take-home:

- Problem solving as search
- Uninformed search
  - Breadth-first search
  - Uniform-cost search
  - Depth-first search
  - Iterative deepening
  - Bidirectional search
- Can you unify them (except bidirectional) using the same algorithm, with different priority functions?
- Performance measures:  
Completeness, optimality, time complexity, space complexity