

Lecture 4:

Advanced Search

CSSE 5600/6600: Artificial Intelligence

Instructor: Bo Liu

[Based on slides from Andrew Moore <http://www.cs.cmu.edu/~awm/tutorials>]


Optimization problems

- Previously we want a **path** from start to goal
 - **Uninformed search**: $g(s)$: Iterative Deepening
 - **Informed search**: $g(s)+h(s)$: A*
- **Now a different setting:**
 - Each state s has a **score** $f(s)$ that we can compute
 - The goal is to find the state with the **highest score**, or a **reasonably high score**
 - Do not care about the path
 - This is an **optimization problem**
 - Enumerating the states is intractable
 - Even previous search algorithms are too expensive

Examples

- N-queen: $f(s)$ = number of conflicting queens in state s


Note we want s with the lowest score $f(s)=0$. The techniques are the same. Low or high should be obvious from context.




Examples

- N-queen: $f(s)$ = number of conflicting queens in state s

Note we want s with the lowest score $f(s)=0$. The techniques are the same. Low or high should be obvious from context.




- Traveling salesperson problem (TSP)
 - Visit each city once, return to first city
 - State = order of cities, $f(s)$ = total mileage




Examples

- N-queen: $f(s)$ = number of conflicting queens in state s

Note we want s with the lowest score $f(s)=0$. The techniques are the same. Low or high should be obvious from context.



- Traveling salesperson problem (TSP)
 - Visit each city once, return to first city
 - State = order of cities, $f(s)$ = total mileage
- Boolean satisfiability (e.g., 3-SAT)
 - State = assignment to variables
 - $f(s)$ = # satisfied clauses



$A \vee \neg B \vee C$
 $\neg A \vee C \vee D$
 $B \vee D \vee \neg E$
 $\neg C \vee \neg D \vee \neg E$
 $\neg A \vee \neg C \vee E$

1. HILL CLIMBING




Hill climbing

- Very simple idea: Start from some state s ,
 - Move to a neighbor t with better score. Repeat.
- **Question:** what's a neighbor?
 - You have to define that!
 - The neighborhood of a state is the set of neighbors
 - Also called 'move set'
 - Similar to successor function


Neighbors: N-queen

- Example: N-queen (one queen per column). One possibility:



s


$$f(s)=1$$



Neighborhood
of s


Neighbors: N-queen

- Example: N-queen (one queen per column). One possibility: tie breaking more promising?
 - Pick the right-most most-conflicting column;
 - Move the queen in that column vertically to a different location.






Neighbors: TSP

- state: A-B-C-D-E-F-G-H-A
- f = length of tour



Neighbors: TSP

- state: A-B-C-D-E-F-G-H-A
- f = length of tour
- One possibility: 2-change



Neighbors: SAT

- State: (A=T, B=F, C=T, D=T, E=T)
- f = number of satisfied clauses
- Neighbor:

$$A \vee \neg B \vee C$$

$$\neg A \vee C \vee D$$

$$B \vee D \vee \neg E$$

$$\neg C \vee \neg D \vee \neg E$$

$$\neg A \vee \neg C \vee E$$

Neighbors: SAT

- State: (A=T, B=F, C=T, D=T, E=T)
- f = number of satisfied clauses
- Neighbor: flip the assignment of one variable

(A=**F**, B=F, C=T, D=T, E=T)

(A=T, B=**T**, C=T, D=T, E=T)

(A=T, B=F, C=**F**, D=T, E=T)

(A=T, B=F, C=T, D=**F**, E=T)

(A=T, B=F, C=T, D=T, E=**F**)

$A \vee \neg B \vee C$

$\neg A \vee C \vee D$

$B \vee D \vee \neg E$

$\neg C \vee \neg D \vee \neg E$

$\neg A \vee \neg C \vee E$

Hill climbing

- **Question:** What's a neighbor?
 - (vaguely) Problems tend to have structures. A small change produces a neighboring state.
 - The neighborhood must be small enough for efficiency
 - Designing the neighborhood is critical. This is the real ingenuity – not the decision to use hill climbing.
- **Question:** Pick which neighbor?
- **Question:** What if no neighbor is better than the current state?

Hill climbing

- **Question:** What's a neighbor?
 - (vaguely) Problems tend to have structures. A small change produces a neighboring state.
 - The neighborhood must be small enough for efficiency
 - Designing the neighborhood is critical. This is the real ingenuity – not the decision to use hill climbing.
- **Question:** Pick which neighbor? The best one (greedy)
- **Question:** What if no neighbor is better than the current state? Stop. (Doh!)

Hill climbing algorithm

1. Pick initial state s
2. Pick t in neighbors(s) with the largest $f(t)$
3. IF $f(t) \leq f(s)$ THEN stop, return s
4. $s = t$. GOTO 2.

- Not the most sophisticated algorithm in the world.
- Very greedy.
- Easily stuck.

Hill climbing algorithm

1. Pick initial state s
2. Pick t in neighbors(s) with the largest $f(t)$
3. IF $f(t) \leq f(s)$ THEN stop, return s
4. $s = t$. GOTO 2.


- Not the most sophisticated algorithm in the world.
- Very greedy.
- Easily stuck.

your enemy:


**local
optima**

Local optima in hill climbing

- Useful conceptual picture: f surface = 'hills' in state space




- But we can't see the landscape all at once. Only see the neighborhood. Climb in fog.




Local optima in hill climbing

- Local optima (there can be many!)



- Plateaux



Local optima in hill climbing

- Local optima (there can be many)

Declare to
the world

The rest of the lecture is
about

**Escaping
local optima**

- Plateaus

Where shall I go?

Not every local minimum should be escaped



Variation 1: hill climbing with random restarts

- Very simple modification
 1. When stuck, pick a random new start, run basic hill climbing from there.
 2. Repeat this k times.
 3. Return the best of the k local optima.
- Can be very effective
- Should be tried whenever hill climbing is used

Variations 2 and 3 of hill climbing

- **Question:** How do we make hill climbing less greedy?
 - Stochastic hill climbing
 - Randomly select among better neighbors
 - The better, the more likely
 - Pros / cons compared with basic hill climbing?

Variations 2 and 3 of hill climbing

- **Question:** How do we make hill climbing less greedy?
 - Stochastic hill climbing
 - Randomly select among better neighbors
 - The better, the more likely
 - Pros / cons compared with basic hill climbing?
- **Question:** What if the neighborhood is too large to enumerate? (e.g. N-queen if we need to pick both the column and the move within it)

Variations 2 and 3 of hill climbing

- **Question:** How do we make hill climbing less greedy?
 - Stochastic hill climbing
 - Randomly select among better neighbors
 - The better, the more likely
 - Pros / cons compared with basic hill climbing?
- **Question:** What if the neighborhood is too large to enumerate? (e.g. N-queen if we need to pick both the column and the move within it)
 - First-choice hill climbing
 - Randomly generate neighbors, one at a time
 - If better, take the move
 - Pros / cons compared with basic hill climbing?

Variation 4 of hill climbing

- We are still greedy! Only willing to move upwards.
- Important observation in life:

Sometimes one
needs to
temporarily step
back in order to
move forward.

=

Sometimes one
needs to move to an
inferior neighbor in
order to escape a
local optimum.

Variation 4 of hill climbing

WALKSAT [Selman]

- Consider 3 neighbors
- If any improves f , accept the best
- If none improves f :
 - 50% of the time pick the least bad neighbor
 - 50% of the time pick a random neighbor

This is the best known algorithm for satisfying Boolean formulae.

$$\begin{array}{l} A \vee \neg B \vee C \\ \neg A \vee C \vee D \\ B \vee D \vee \neg E \\ \neg C \vee \neg D \vee \neg E \\ \neg A \vee \neg C \vee E \end{array}$$